


Can Pretrained CLTs Distinguish Hint-Reading from Format Changes?

The hint-specific signal (bottom) is ~10x smaller than the total GRPO change (top)

With-Hint and No-Hint Differentials Are Nearly Identical



Isolating the Hint-Reading Signal: Very Small Compared to Format Changes

